

Chapter 8: Convolutional Networks

Dietrich Klakow

Spoken Language Systems

Saarland University, Germany

`Dietrich.Klakow@LSV.Uni-Saarland.De`

Neural Networks Implementation and Application




UNIVERSITÄT
DES
SAARLANDES



Source: <http://www.dailymail.co.uk/news/article-564264/The-worlds-tallest-Lego-tower-took-500-000-bricks-build.html>

Introduction: AlexNet



- ▶ Source: "ImageNet Classification with Deep Convolutional Neural Networks"
Alex Krizhevsky, Ilya Sutskever, Geoffrey E Hinton, NIPS 2012
- ▶ 650,000 neurons
- ▶ 60 million parameters
- ▶ Roughly 100 parameters per neuron!



- ① Convolution
- ② Motivation for Convolution
- ③ Pooling
- ④ Variants of the convolution function
- ⑤ Structured outputs and other related issues
- ⑥ Example: cifar10 from tensorflow tutorial

Convolution



- ▶ Operates on two functions
- ▶ Example signal $s(t)$ that you are hearing depends on
 - ▶ Signal produced $x(t)$
 - ▶ Impulse response of lecture hall $w(t)$

▶

$$s(t) = \int x(a)w(t-a)da$$

- ▶ Shorthand notation

$$s(t) = (x * w)(t)$$

- ▶ When $w(t)$ is sometimes called *kernel*
- ▶ Convolution is commutative



- ▶ Discrete convolution (time index t)

$$s(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a)$$

- ▶ Two dimensional discrete convolution

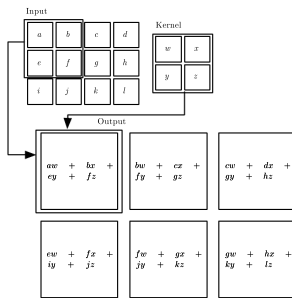
$$S(i,j) = (I * K)(i,j) = \sum_m \sum_n I(i-m, j-n)K(m, n)$$

- ▶ Relation to *cross-correlation*

$$S(i,j) = (I * K)(i,j) = \sum_m \sum_n I(i+m, j+n)K(m, n)$$

- ▶ Cross-correlation is convolution with a flipped kernel
- ▶ For this reason the book uses cross-correlation and convolution interchangeably

Example of a 2D convolution




- Strictly speaking this calculates a correlation

Motivation for Convolution




- ▶ Parameter sharing in a layer of a neural network
- ▶ Provides means to work with inputs of variable size
- ▶ Faster computation

Sparse connectivity viewed from below




- x_3 influences only s_2, s_3 and s_4

Sparse connectivity viewed from above




- ▶ s_3 is only influenced by x_2 , x_3 and x_4
- ▶ The neurons that influence s_3 are called the *receptive field* of s_3

Receptive field in a deeper network



- ▶ Locally the connections are sparse
- ▶ Globally, everything indirectly influences most (or all) of the output




- ▶ Take difference between each pixel and its vertical neighbor
- ▶ Image is 280x320 pixels
- ▶ This edge detection takes 267960 operations
- ▶ Applying a fully connected matrix would take 16 billion floating point operations

Pooling



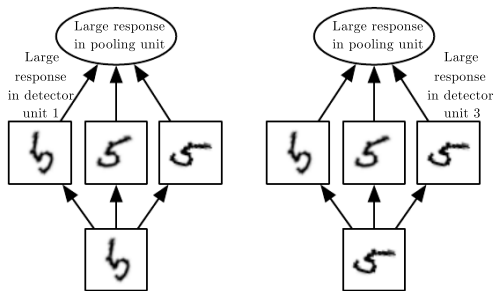
- ▶ *Pooling*: replaces the output at a certain location by some summary statistics over neighborhood
- ▶ *max pooling*: maximum output within a regular window
- ▶ Alternative: average over a rectangular window

Example of max pooling



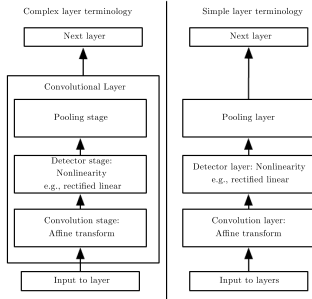
- ▶ Input in the two examples above is shifted by a pixel
- ▶ Output becomes invariant so small shifts

Example of learned invariances




- Orientation of the input does not matter for the output

Typical layout of a convolutional network



- Alternative terminologies exist as to what is a layer

Pooling with down-sampling



- Reduces computational burden for next layer

Variants of the convolution function

Convolution for multi-channel input (color images)



- ▶ Let $V_{l,j,k}$ be a color image where l denotes the color channel and j and k the position in the image

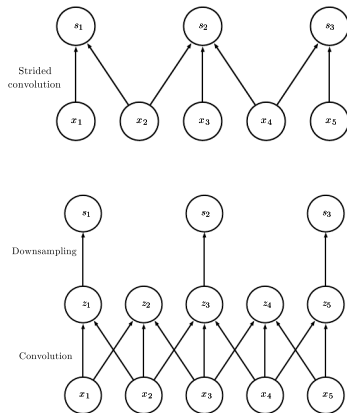


$$Z_{i,j,k} = \sum_{l,m,n} V_{l,j+m,k+n} K_{i,l,m,n}$$


- ▶ Above formula assumes C/Python indexing in the arrays
- ▶ *stride*: skip some of the outputs
- ▶ Convolution with stride

$$Z_{i,j,k} = \sum_{l,m,n} V_{l,(j-1) \times s + m, (k-1) \times s + n} K_{i,l,m,n}$$

- ▶ s is the stride
- ▶ Indices i and l are indexing color
- ▶ j, k, m, n : space indices



- Convolution with stride:
equivalent to convolution with down-sampling
- Issue: s_1 and s_3 are influenced only by two inputs



- Pooling results in an undesired reduction in the number of nodes



- Zeros (black circles) are added to keep layer size fixed




$$Z_{i,j,k} = \sum_{l,m,n} V_{l,j+m,k+n} K_{i,j,k,l,m,n}$$


Use when:

- ▶ all interactions are local
- ▶ different regions of image behave differently
- ▶ Indices i and l are indexing color
- ▶ j, k, m, n : space indices

Comparison of full weights and convolution




- ▶ Top: locally connected
- ▶ Center: convolution
- ▶ Bottom: fully connected



- Compromise between locally connected and convolution (see next slide)
- No connections between different channels (e.g. colors)
- Formal definition

$$Z_{i,j,k} = \sum_{l,m,n} V_{l,j+m,k+n} K_{i,l,m,n,j\%t+1,k\%t+1}$$

- % is the modulo operation



- ▶ Top: general locally connected network
- ▶ Center: $t=4$
- ▶ Bottom: $t=2$



- ▶ Suppose we want to minimize $J(V, K)$
- ▶ Suppose during back propagation we receive tensor

$$G_{i,j,k} = \frac{\partial}{\partial Z_{i,j,k}} J(V, K)$$

- ▶ Gradient to update K


$$g(G, V, s)_{i,j,k,l} = \frac{\partial}{\partial K_{i,j,k,l}} J(V, K) = \sum_{m,n} G_{i,m,n} V_{j,(m-1) \times s + k, (n-1) \times s + l}$$

- ▶ To backpropagate further to the beginning of the network

$$h(K, V, s)_{i,j,k} = \frac{\partial}{\partial V_{i,j,k}} J(V, K) = \sum_{\substack{m, n, s. t. \\ (l-1) \times s + m = j}} \sum_{\substack{n, p, s. t. \\ (n-1) \times s + p = k}} \sum_q K_{q,i,m,p} G_{q,l,n}$$


- ▶ Note: forget about the details! Not relevant for exam in this generality.

Structured outputs and other related issues



Source: <https://www.mathworks.com/matlabcentral/mlc-downloads/downloads/submissions/29517/versions/5/screenshot.jpg>

In image segmentation an image X produces an output \hat{Y} for each pixel



- ▶ X is the matrix representing the input image
- ▶ $\hat{Y}^{(i)}$ is the matrix with labels for each input pixel
- ▶ Refining output for image segmentation iteratively
- ▶ Simple recurrent network
- ▶ Captures interaction between neighboring pixels



- ▶ Single channel
 - ▶ 1-D: audio waveform
 - ▶ 2-D: grey level image
 - ▶ 3-D: CT scan
- ▶ Multi-channel
 - ▶ 1-D: Movement of skeleton (one channel per joint)
 - ▶ 2-D: Color image
 - ▶ 3-D: Color video



Assume kernel is w pixels wide

Kernel is d dimensional

→ $O(w^d)$ operations

- ▶ Fourier transform
 - ▶ Transform kernel and image to frequency space: $O(d \times w \log w)$ operations
 - ▶ Multiply in frequency space: $O(d \times w)$ operations
 - ▶ Transform back to position space: $O(d \times w \log w)$ operations
- ▶ Separable kernels
 - ▶ Kernel can be written as an out product of vectors
 - ▶ That means that the kernel can be applied in each dimension separately
 - ▶ Requires $O(d \times w)$ operations
 - ▶ Note: only some kernels are separable



- ▶ Let $I(x, y)$ be an image
- ▶ and $w(x, y)$ a filter
- ▶ Filtered image is given by

$$S(x_f, y_f) = \sum_{x, y} w(x, y) I(x - x_f, y - y_f)$$

- ▶ Gabor function

$$w(x, y) = \alpha \exp(-\beta_x x'^2 - \beta_y y'^2) \cos(fx' + \phi)$$

- ▶ where




$$x' = (x - x_0) \cos(\tau) + (y - y_0) \sin(\tau)$$

- ▶ and

$$y' = -(x - x_0) \sin(\tau) + (y - y_0) \cos(\tau)$$


- ▶ $\beta_x, \beta_y, f, x_0, y_0, \phi, \tau$ are parameters

Gabor functions for a variety of parameters settings



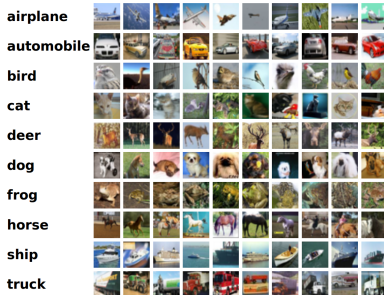
- ▶ Gabor functions probe texture on different directions with different “wave length”

Features learned in first layer of CNN



- Resembles Gabor functions

Example: cifar10 from tensorflow tutorial



- ▶ 60000 32x32 color images
- ▶ 10 classes,
- ▶ 6000 images per class
- ▶ 50000 training images
- ▶ 10000 test images.

Architecture of the model



2-D convolution given 4-D input



```
tf.nn.conv2d(input, filter, strides, padding,  
use_cudnn_on_gpu=None, data_format=None, name=None)
```

Given:

- ▶ an input tensor of shape
[batch, in_height, in_width, in_channels]
- ▶ A filter / kernel tensor of shape
[filter_height, filter_width, in_channels, out_channels]

In detail, with the default NHWC format,

```
output[b, i, j, k] =  
    sum_di, dj, q input[b, strides[1] * i + di,  
        strides[2] * j + dj, q] * filter[di, dj, q, k]
```



```
tf.nn.max_pool_with_argmax(input, ksize, strides, padding,  
Targmax=None, name=None)
```

Performs max pooling on the input and outputs both max values and indices.

The indices in argmax are flattened, so that a maximum value at position $[b, y, x, c]$ becomes flattened index $((b * \text{height} + y) * \text{width} + x) * \text{channels} + c$.



```
tf.nn.local_response_normalization(input, depth_radius=None,  
bias=None, alpha=None, beta=None, name=None)
```

- ▶ The 4-D input tensor is treated as a 3-D array of 1-D vectors (along the last dimension)
- ▶ each vector is normalized independently
- ▶ Within a given vector, each component is divided by the weighted, squared sum of inputs within `depth_radius`.

In detail:


```
sqr_sum[a, b, c, d] =  
    sum(input[a, b, c, d-depth_radius : d+depth_radius+1]**2)  
output = input / (bias + alpha * sqr_sum) ** beta
```

Purpose: “lateral inhibition” that is an excited neuron reduces the neighbors in importance


Example: cifar10




Walk through cifar10.py starting line 204



Learning rate





- ▶ Calculate gradients for different batches on different GPUs
- ▶ Assumes that all GPUs need the same time for any batch



- ▶ Convolutional networks root in image processing
- ▶ Together with pooling they are good to cover invariances
- ▶ Surprisingly they work for language as well